
PROGRAM 1: Linear Regression for the diabetes dataset

REGNO:24251115

DATE: 23-05-2025

```
from sklearn.datasets import load_diabetes
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score

# Load the dataset
diabetes = load_diabetes()
X = diabetes.data
y = diabetes.target

# Split into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Create and train the model
model = LinearRegression()
model.fit(X_train, y_train)

# Make predictions
y_pred = model.predict(X_test)
comparison_df = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})
comparison_df

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
```

```
r2 = r2_score(y_test, y_pred)
```

```
# Output results
```

```
print("Coefficients:", model.coef_)
```

```
print("Intercept:", model.intercept_)
```

```
print("Mean Squared Error:", mse)
```

```
print("R2 Score:", r2)
```

OUTPUT:

```
model = LinearRegression()
```

```
model
```

```
LinearRegression
```

	Actual	Predicted
304	253.0	117.614637
358	90.0	48.721858
128	115.0	90.586456
373	168.0	143.619232
9	310.0	212.757922
...
129	268.0	214.506193
63	128.0	96.055855
342	178.0	159.934376
291	248.0	194.817561
329	135.0	104.155230

89 rows × 2 columns

```
print("Coefficients:", model.coef_)
```

```
Coefficients: [ 37.90402135 -241.96436231 542.42875852 347.703
84391 -931.48884588
518.06227698 163.41998299 275.31790158 736.1988589 48.670
65743]
```

```
print("Intercept:", model.intercept_)
```

```
Intercept: 151.34560453985995
```

```
print("Mean Squared Error:", mse)
```

```
Mean Squared Error: 2900.1936284934804
```

```
print("R2 Score:", r2)
```

```
R2 Score: 0.4526027629719196
```

PROGRAM 2: Logistic Regression for the breast cancer dataset

REGNO: 24251115

DATE: 23-05-2025

```
from sklearn.datasets import load_breast_cancer

from sklearn.linear_model import LogisticRegression

from sklearn.model_selection import train_test_split

from sklearn.metrics import classification_report, confusion_matrix, accuracy_score


# Load the dataset

cancer = load_breast_cancer()

X = cancer.data

y = cancer.target


# Split into train and test sets

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


# Create and train the model

model = LogisticRegression(max_iter=10000) # max_iter increased to ensure convergence

model.fit(X_train, y_train)


# Make predictions

y_pred = model.predict(X_test)

comparison_df = pd.DataFrame(

    {

        'Actual': y_test,

        'Predicted': y_pred

    }

)
```

```
comparison_df
```

```
# Evaluate the model
```

```
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
```

```
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

```
print("Accuracy Score:", accuracy_score(y_test, y_pred))
```

OUTPUT:

```
LogisticRegression
LogisticRegression(max_iter=10000)

model.fit(X_train, y_train)

LogisticRegression
LogisticRegression(max_iter=10000)
```

	Actual	Predicted
0	1	1
1	0	0
2	0	0
3	1	1
4	1	1
...
109	1	1
110	0	0
111	1	1
112	1	0
113	0	0

114 rows × 2 columns

Confusion Matrix:

```
[[39  4]
 [ 1 70]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.97	0.91	0.94	43
1	0.95	0.99	0.97	71
accuracy			0.96	114
macro avg	0.96	0.95	0.95	114
weighted avg	0.96	0.96	0.96	114

Accuracy Score: 0.956140350877193

PROGRAM 3: Classification matrix for the Decision Tree classifier for iris dataset

REGNO:24251115

DATE: : 23-05-2025

```
sklearn.datasets import load_iris

from sklearn.tree import DecisionTreeClassifier

from sklearn.model_selection import train_test_split

from sklearn.metrics import confusion_matrix, classification_report, accuracy_score


# Load the iris dataset

iris = load_iris()

X = iris.data

y = iris.target


# Split into train and test sets

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


# Create and train Decision Tree classifier

clf = DecisionTreeClassifier(random_state=42)

clf.fit(X_train, y_train)


# Predict on test data

y_pred = clf.predict(X_test)


# Evaluate using confusion matrix

cm = confusion_matrix(y_test, y_pred)

report = classification_report(y_test, y_pred, target_names=iris.target_names)

accuracy = accuracy_score(y_test, y_pred)
```

Output

```
print("Confusion Matrix:\n", cm)
print("\nClassification Report:\n", report)
print("Accuracy Score:", accuracy)
```

OUTPUT:

```
clf = DecisionTreeClassifier(random_state=42)
```

```
clf
```

```
▼      DecisionTreeClassifier      ⓘ ?
DecisionTreeClassifier(random_state=42)
```

Confusion Matrix:

```
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
```

Classification Report:

	precision	recall	f1-score	support
setosa	1.00	1.00	1.00	10
versicolor	1.00	1.00	1.00	9
virginica	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

Accuracy Score: 1.0

PROGRAM 4: Classification model for Wine dataset using RandomForestClassifier

REGNO: 24251115

DATE: : 23-05-2025

```
from sklearn.datasets import load_wine
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Load the dataset
wine = load_wine()
X = wine.data
y = wine.target

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize the RandomForestClassifier
clf = RandomForestClassifier(n_estimators=100, random_state=42)

# Train the model
clf.fit(X_train, y_train)

# Make predictions
y_pred = clf.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
report = classification_report(y_test, y_pred, target_names=wine.target_names)
```

```
conf_matrix = confusion_matrix(y_test, y_pred)
```

```
# Display results
```

```
print("Accuracy:", accuracy)
```

```
print("Classification Report:\n", report)
```

```
print("Confusion Matrix:\n", conf_matrix)
```

OUTPUT:

```
clf = RandomForestClassifier(n_estimators=100, random_state=42)
```

```
clf
```

▼ RandomForestClassifier ⓘ ?

```
RandomForestClassifier(random_state=42)
```

Accuracy: 1.0

Classification Report:

	precision	recall	f1-score	support
class_0	1.00	1.00	1.00	14
class_1	1.00	1.00	1.00	14
class_2	1.00	1.00	1.00	8
accuracy			1.00	36
macro avg	1.00	1.00	1.00	36
weighted avg	1.00	1.00	1.00	36

Confusion Matrix:

```
[[14  0  0]
 [ 0 14  0]
 [ 0  0  8]]
```

PROGRAM 5: Classification model for Wine dataset using K-Nearest Neighbors Classifier.

REGNO: 24251115

DATE: : 23-05-2025

Import necessary libraries

from sklearn.datasets import load_wine

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler

from sklearn.neighbors import KNeighborsClassifier

from sklearn.metrics import classification_report, accuracy_score, confusion_matrix

Step 1: Load the dataset

wine = load_wine()

X = wine.data

y = wine.target

Step 2: Split the dataset into training and testing sets

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

Step 3: Standardize the features

scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train)

X_test_scaled = scaler.transform(X_test)

Step 4: Create and train the KNN model

knn = KNeighborsClassifier(n_neighbors=5)

knn.fit(X_train_scaled, y_train)

```
# Step 5: Make predictions
```

```
y_pred = knn.predict(X_test_scaled)
```

```
# Step 6: Evaluate the model
```

```
print("Confusion Matrix:")
```

```
print(confusion_matrix(y_test, y_pred))
```

```
print("\nClassification Report:")
```

```
print(classification_report(y_test, y_pred))
```

```
print("Accuracy Score:", accuracy_score(y_test, y_pred))
```

OUTPUT:

```
knn = KNeighborsClassifier(n_neighbors=5)
```

```
knn
```

▼ KNeighborsClassifier ⓘ ?

```
KNeighborsClassifier()
```

Confusion Matrix:

```
[[14  0  0]
 [ 1 12  1]
 [ 0  0  8]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.93	1.00	0.97	14
1	1.00	0.86	0.92	14
2	0.89	1.00	0.94	8
accuracy			0.94	36
macro avg	0.94	0.95	0.94	36
weighted avg	0.95	0.94	0.94	36

Accuracy Score: 0.9444444444444444

PROGRAM 6: Regression model for Fetch House Pricing using RandomForestRegressor

REGNO: 24251115

DATE: : 23-05-2025

```
import matplotlib.pyplot as plt

from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score


# Step 1: Load the dataset
housing = fetch_california_housing()
X = housing.data
y = housing.target


# Step 2: Split into training and testing data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


# Step 3: Standardize the data
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)


# Step 4: Train the Random Forest Regressor
model = RandomForestRegressor(n_estimators=100, random_state=42)
model.fit(X_train_scaled, y_train)
```

```
# Step 5: Predict
```

```
y_pred = model.predict(X_test_scaled)
```

```
# Step 6: Evaluate the model
```

```
mse = mean_squared_error(y_test, y_pred)
```

```
r2 = r2_score(y_test, y_pred)
```

```
print("Mean Squared Error:", round(mse, 4))
```

```
print("R-squared Score:", round(r2, 4))
```

```
# Step 7: Plot actual vs predicted values
```

```
plt.scatter(y_test, y_pred, alpha=0.5, color='green')
```

```
plt.xlabel("Actual House Prices")
```

```
plt.ylabel("Predicted House Prices")
```

```
plt.title("Actual vs. Predicted House Prices")
```

```
plt.grid(True)
```


```
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], 'r--') # Diagonal line
```

```
plt.show()
```

OUTPUT:

Mean Squared Error: 0.2552

R-squared Score: 0.8053



PROGRAM 7: Regression model for Fetch House Prizing using KNeighborRegressor

REGNO: 24251115

DATE: 23-05-2025

```
import matplotlib.pyplot as plt
from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import mean_squared_error, r2_score

# Step 1: Load the dataset
housing = fetch_california_housing()
X = housing.data
y = housing.target

# Step 2: Split into training and testing data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Step 3: Standardize the data
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Step 4: Train the K-Nearest Neighbors Regressor
model = KNeighborsRegressor(n_neighbors=5)
model.fit(X_train_scaled, y_train)

# Step 5: Predict
y_pred = model.predict(X_test_scaled)
```

```
# Step 6: Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)


print("Mean Squared Error:", round(mse, 4))
print("R-squared Score:", round(r2, 4))

# Step 7: Plot actual vs predicted values
plt.scatter(y_test, y_pred, alpha=0.5, color='blue')
plt.xlabel("Actual House Prices")
plt.ylabel("Predicted House Prices")
plt.title("Actual vs. Predicted House Prices (KNN)")
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], 'r--') # Diagonal line
plt.grid(True)
plt.show()
```

OUTPUT:

Mean Squared Error: 0.4324

R-squared Score: 0.67



PROGRAM 8: Preprocessing missing value treatment

REGNO: 24251115

DATE: 23-05-2025

```
import matplotlib.pyplot 8. Preprocessing missing value treatment
```

```
import pandas as pd
```

```
from sklearn.impute import SimpleImputer
```

```
from sklearn.linear_model import LinearRegression
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.metrics import mean_squared_error
```

```
# Sample data with missing values
```

```
data = {
```

```
    'Age': [25, 30, None, 35, 40],
```

```
    'Salary': [50000, 60000, 52000, None, 58000],
```

```
    'HousePrice': [200000, 250000, 220000, 240000, 230000]
```

```
}
```

```
df = pd.DataFrame(data)
```

```
# Separate features and target
```

```
X = df[['Age', 'Salary']]
```

```
y = df['HousePrice']
```

```
# Fill missing values with mean
```

```
imputer = SimpleImputer(strategy='mean')
```

```
X = imputer.fit_transform(X)
```

```
# Train-test split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Train linear regression model
model = LinearRegression()
model.fit(X_train, y_train)

# Predict and evaluate
y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)

print("Mean Squared Error:", round(mse, 2))
```

OUTPUT:

```
model
```

▼ LinearRegression ⓘ ?

```
LinearRegression()
```

```
X = df[['Age', 'Salary']]
```

```
X
```

	Age	Salary
0	25.0	50000.0
1	30.0	60000.0
2	NaN	52000.0
3	35.0	NaN
4	40.0	58000.0

```
y = df['HousePrice']
```

```
y
```

0	200000
1	250000
2	220000
3	240000
4	230000

Name: HousePrice, dtype: int64

PROGRAM 9: Preprocessing, Scaling and Encoding

REGNO: 24251115

DATE: 23-05-2025

```
import pandas as pd

from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score


data = {
    'Age': [25, 30, None, 35, 40],
    'Salary': [50000, 60000, 52000, None, 58000],
    'City': ['New York', 'Paris', 'London', None, 'Paris'],
    'HousePrice': [200000, 250000, 220000, 240000, 230000]
}

df = pd.DataFrame(data)
X = df.drop('HousePrice', axis=1)
y = df['HousePrice']

# Missing value treatment
num_cols = ['Age', 'Salary']
cat_cols = ['City']

imputer_num = SimpleImputer(strategy='mean')
X[num_cols] = imputer_num.fit_transform(X[num_cols])
```

```
imputer_cat = SimpleImputer(strategy='most_frequent')
X[cat_cols] = imputer_cat.fit_transform(X[cat_cols])

# Encoding categorical variables
X = pd.get_dummies(X, drop_first=True)

# Scaling numeric features
scaler = StandardScaler()
X[num_cols] = scaler.fit_transform(X[num_cols])

# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train model
model = LinearRegression()
model.fit(X_train, y_train)

# Predict and evaluate
y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Mean Squared Error: {mse:.2f}")
print(f"R-squared Score: {r2:.2f}")
```


OUTPUT

df

	Age	Salary	City	HousePrice
0	25.0	50000.0	New York	200000
1	30.0	60000.0	Paris	250000
2	NaN	52000.0	London	220000
3	35.0	NaN	None	240000
4	40.0	58000.0	Paris	230000

model

LinearRegression ⓘ ?

LinearRegression()

y_pred

array([237174.72118959])

mse

164487776.5647242

r2

nan

PROGRAM 10: SVM model with preprocessing for the Gender classification dataset

REGNO: 24251115

DATE: 23-05-2025

```
import pandas as pd

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.svm import SVC
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
import matplotlib.pyplot as plt
import seaborn as sns

# Sample dataset
data = {
    'Height': [170, 160, 180, 175, 155, 165, 172, 158],
    'Weight': [65, 55, 80, 70, 50, 60, 68, 54],
    'Age': [25, 22, 28, 30, 20, 27, 24, 21],
    'Gender': ['Male', 'Female', 'Male', 'Male', 'Female', 'Female', 'Male', 'Female']
}

df = pd.DataFrame(data)

# Features and target
X = df.drop('Gender', axis=1)
y = df['Gender']

# Encode target
label_encoder = LabelEncoder()
y_encoded = label_encoder.fit_transform(y) # Male=1, Female=0 or vice versa
```

```
# Scale features

scaler = StandardScaler()

X_scaled = scaler.fit_transform(X)


# Split data

X_train, X_test, y_train, y_test = train_test_split(X_scaled, y_encoded, test_size=0.25,
random_state=42)


# Train SVM

svm_model = SVC(kernel='linear', random_state=42)

svm_model.fit(X_train, y_train)


# Predict

y_pred = svm_model.predict(X_test)


# Evaluate

acc = accuracy_score(y_test, y_pred)

print(f'Accuracy: {acc:.2f} ")

print(classification_report(y_test, y_pred, target_names=label_encoder.classes_))


# Confusion Matrix

cm = confusion_matrix(y_test, y_pred)


# Plot confusion matrix

plt.figure(figsize=(6,4))

sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',

            xticklabels=label_encoder.classes_,

            yticklabels=label_encoder.classes_)


plt.xlabel('Predicted')
```

```
plt.ylabel('Actual')
plt.title('Confusion Matrix - Gender Classification')
plt.show()
```

OUTPUT:

Accuracy: 0.50

	precision	recall	f1-score	support
Female	1.00	0.50	0.67	2
Male	0.00	0.00	0.00	0
accuracy			0.50	2
macro avg	0.50	0.25	0.33	2
weighted avg	1.00	0.50	0.67	2



PROGRAM 11: Gradient boosting classification with preprocessing for the Student Marks Dataset

REGNO: 24251115

DATE: : 23-05-2025

```
import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.impute import SimpleImputer

from sklearn.preprocessing import StandardScaler, LabelEncoder

from sklearn.ensemble import GradientBoostingClassifier

from sklearn.metrics import accuracy_score, confusion_matrix

import matplotlib.pyplot as plt

import seaborn as sns
```

1. Sample data

```
data = {
    'Math': [85, 70, 90, None, 60],
    'Science': [90, 65, None, 80, 70],
    'English': [78, 80, 85, 75, None],
    'Gender': ['Male', 'Female', 'Male', 'Female', 'Female'],
    'Passed': ['Yes', 'No', 'Yes', 'No', 'No']
}
```

```
df = pd.DataFrame(data)
```

2. Separate features and target

```
X = df.drop('Passed', axis=1)
y = df['Passed']
```

3. Fill missing numeric values with mean

```
imputer = SimpleImputer(strategy='mean')
```

```
X[['Math', 'Science', 'English']] = imputer.fit_transform(X[['Math', 'Science', 'English']])
```

4. Encode Gender column

```
X = pd.get_dummies(X, columns=['Gender'], drop_first=True)
```

5. Scale numeric columns

```
scaler = StandardScaler()
```

```
X[['Math', 'Science', 'English']] = scaler.fit_transform(X[['Math', 'Science', 'English']])
```

6. Encode target variable

```
le = LabelEncoder()
```

```
y = le.fit_transform(y) # Yes=1, No=0
```

7. Split dataset

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

8. Train Gradient Boosting Classifier

```
model = GradientBoostingClassifier(random_state=42)
```

```
model.fit(X_train, y_train)
```

9. Predict and evaluate

```
y_pred = model.predict(X_test)
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
print("Accuracy:", round(accuracy, 2))
```

10. Confusion matrix

```
cm = confusion_matrix(y_test, y_pred)
```

```
# 11. Plot confusion matrix

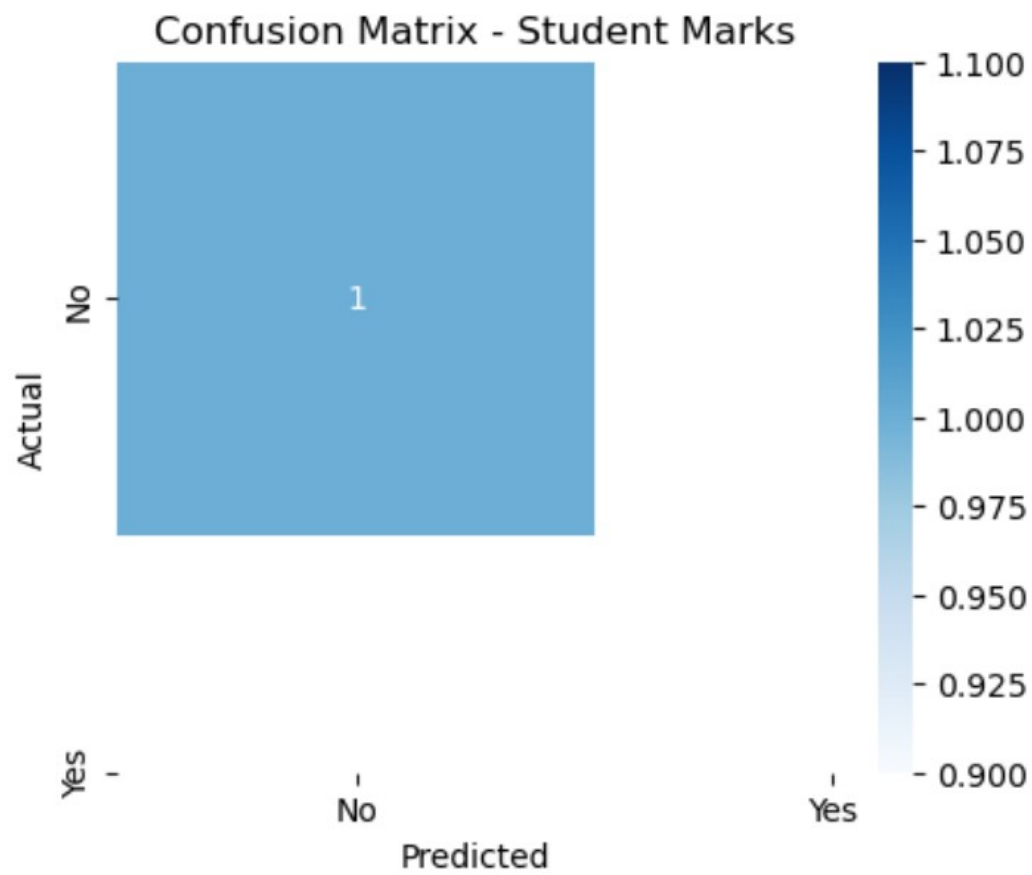
plt.figure(figsize=(5,4))

sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=le.classes_, yticklabels=le.classes_)

plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix - Student Marks')

plt.show()
```

OUTPUT:



```
] : accuracy
```

```
] : 1.0
```

PROGRAM 12: . Visualization using matplotlib for iris dataset

- **Bar chart**
- **Pie chart**
- **Histogram**

REGNO: 24251115

DATE: : 23-05-2025

```
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris

# Load Iris dataset
iris = load_iris()
target = iris.target
target_names = iris.target_names


# Count samples per species
counts = [sum(target == i) for i in range(len(target_names))]

# 1. Bar chart
plt.bar(target_names, counts)
plt.title('Samples per Iris Species')
plt.xlabel('Species')
plt.ylabel('Count')
plt.show()

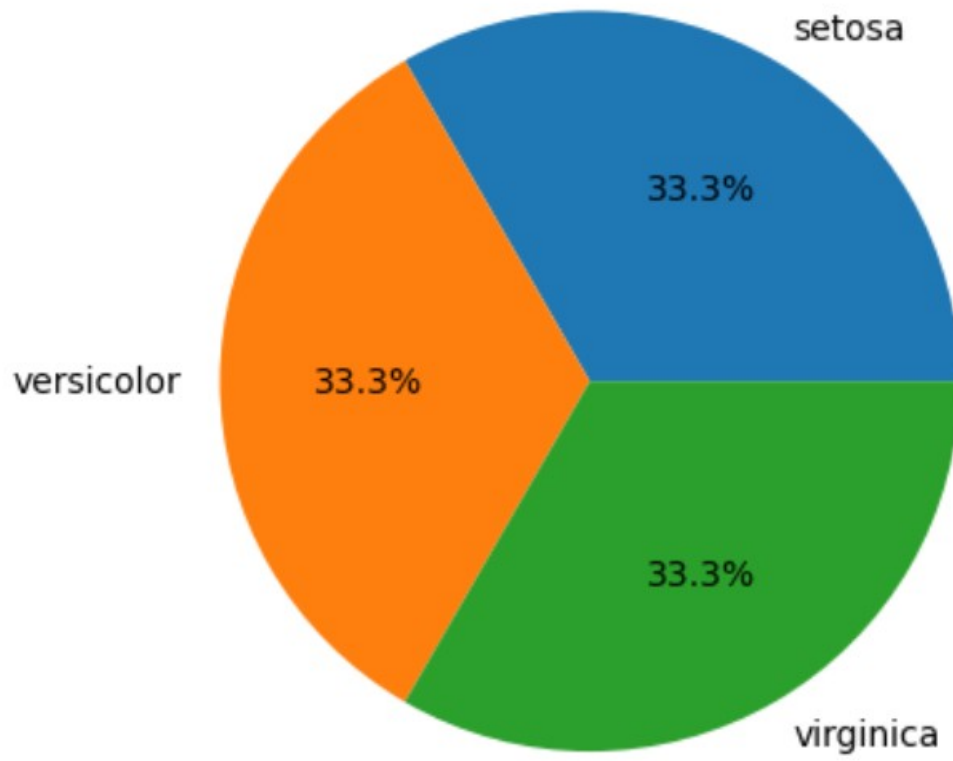
# 2. Pie chart
plt.pie(counts, labels=target_names, autopct='%1.1f%%')
plt.title('Species Distribution')
plt.show()
```


```
# 3. Histogram (Sepal Length)
plt.hist(iris.data[:, 0], bins=15)
plt.title('Sepal Length Distribution')
plt.xlabel('Sepal Length (cm)')
plt.ylabel('Frequency')
plt.show()
```

OUTPUT:



Species Distribution





PROGRAM 13: Pickle File for digits dataset with GaussianNB model

REGNO: 24251115

DATE: 23-05-2025

```
from sklearn.datasets import load_digits
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, confusion_matrix, ConfusionMatrixDisplay
import matplotlib.pyplot as plt
```

1. Load digits dataset

```
digits = load_digits()
```

```
X = digits.data
```

```
y = digits.target
```

2. Split data

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

3. Train Gaussian Naive Bayes model

```
model = GaussianNB()
```

```
model.fit(X_train, y_train)
```

4. Predict and evaluate

```
y_pred = model.predict(X_test)
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
print(f"Accuracy: {accuracy:.2f}")
```

5. Confusion matrix plot

```
cm = confusion_matrix(y_test, y_pred)
```

```
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=digits.target_names)
disp.plot(cmap=plt.cm.Blues)
plt.title("Confusion Matrix - GaussianNB on Digits Dataset")
plt.show()
```

6. Save model to pickle file

```
with open('gaussian_nb_digits.pkl', 'wb') as f:
    pickle.dump(model, f)
print("Model saved to 'gaussian_nb_digits.pkl'")
```

7. Load the model (example)

```
with open('gaussian_nb_digits.pkl', 'rb') as f:
    loaded_model = pickle.load(f)
```

8. Predict using loaded model (example)

```
sample_preds = loaded_model.predict(X_test[:5])
print("Sample predictions from loaded model:", sample_preds)
```


OUTPUT:

```
model
```

▼ GaussianNB ⓘ ?

```
GaussianNB()
```

Accuracy: 0.85



Model saved to 'gaussian_nb_digits.pkl'

Sample predictions from loaded model: [6 9 3 7 2]

PROGRAM 14: Advanced visualization and interpretation using seaborn for titanic dataset

- **Stacked bar chart**
- **KDE plot**
- **Violin plot**
- **Heatmap**
- **Swarm plot**

REGNO: 24251115

DATE: 23-05-2025

Import matplotlib.pyplot as plt

Load dataset

titanic = sns.load_dataset('titanic')

1. Stacked Bar Chart (using pandas plot)

survival = titanic.groupby(['class', 'sex'])['survived'].mean().unstack()

survival.plot(kind='bar', stacked=True, color=['pink', 'lightblue'])

plt.title('Survival Rate by Class and Sex')

plt.ylabel('Survival Rate')

plt.show()

2. KDE Plot for Age by Survival

sns.kdeplot(data=titanic, x='age', hue='survived', fill=True)

plt.title('Age Distribution by Survival')

plt.show()

3. Violin Plot for Age by Class

sns.violinplot(x='class', y='age', data=titanic)


```
plt.title('Age Distribution by Class')
plt.show()
```

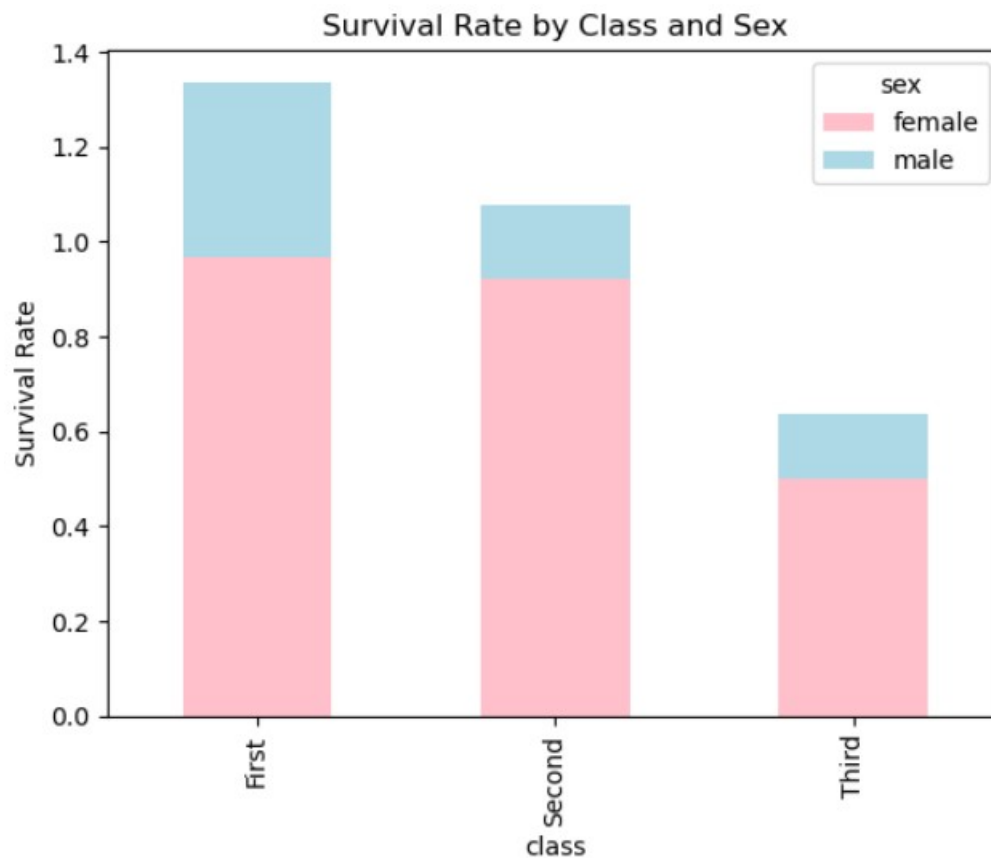
4. Heatmap for Correlation


```
numeric_data = titanic.select_dtypes(include=['float64', 'int64'])
sns.heatmap(numeric_data.corr(), annot=True, cmap='coolwarm')
plt.title('Correlation Heatmap')
plt.show()
```

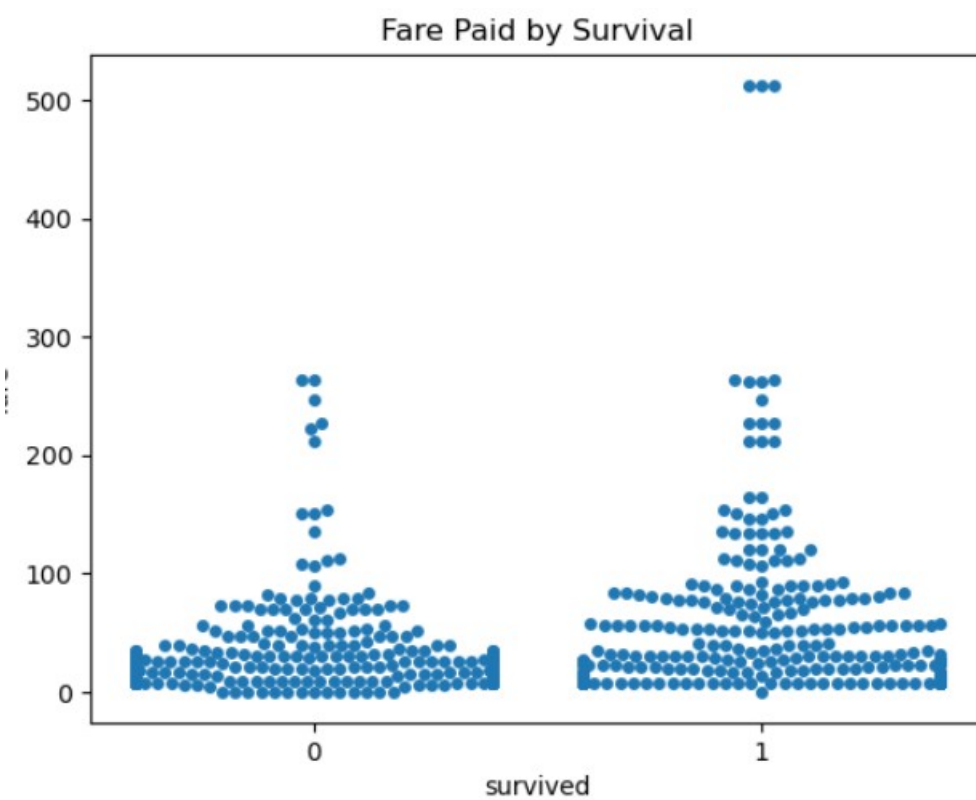
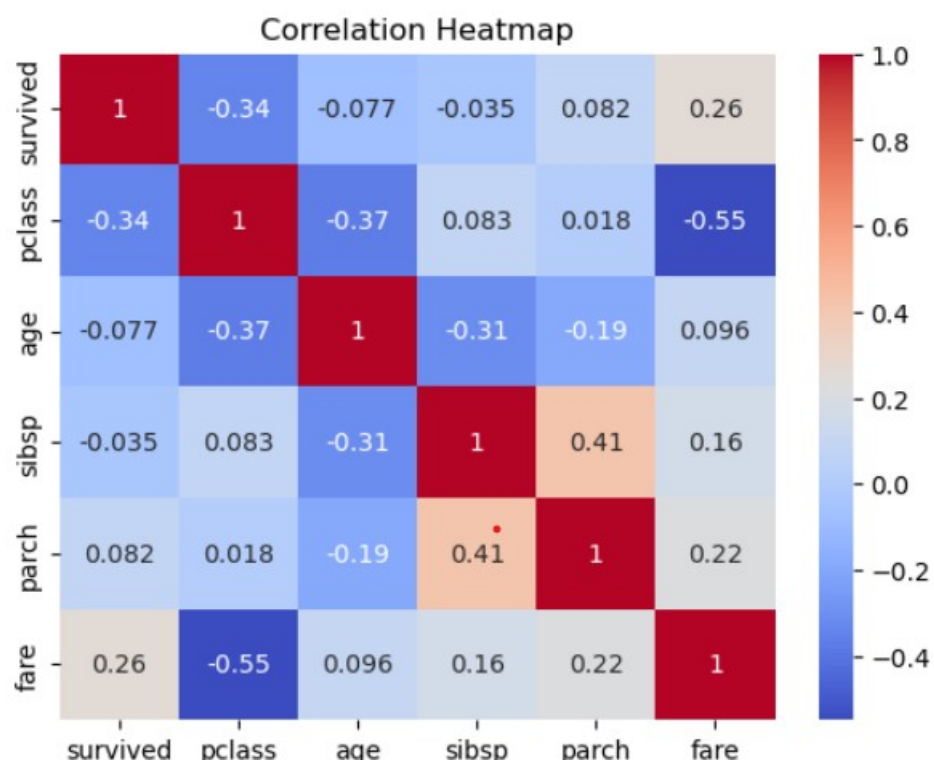
5. Swarm Plot for Fare by Survival

```
sns.swarmplot(x='survived', y='fare', data=titanic)
plt.title('Fare Paid by Survival')
plt.show()
```

OUTPUT:







PROGRAM 15: KMeans Cluster analysis

REGNO: 24251115

DATE: 23-05-2025

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
from sklearn.datasets import load_iris

# Load Iris dataset
iris = load_iris()
df = pd.DataFrame(iris.data, columns=iris.feature_names)

# Apply KMeans clustering
kmeans = KMeans(n_clusters=3, random_state=0)
df['cluster'] = kmeans.fit_predict(df)


# Add species for comparison (optional)
df['species'] = iris.target

# Plot clusters (using first two features for simplicity)
plt.figure(figsize=(7,5))
sns.scatterplot(x=df.iloc[:, 0], y=df.iloc[:, 1], hue=df['cluster'], palette='Set1')
plt.title('KMeans Clustering on Iris Dataset')
plt.xlabel(iris.feature_names[0])
plt.ylabel(iris.feature_names[1])
plt.show()
```

OUTPUT:

```
KMeans
```

```
KMeans(n_clusters=3, random_state=0)
```



PROGRAM 16: . Program to build a machine learning model to predict whether the person is diabetic or not using Bagging Classifier

REGNO:24251115

DATE: 23-05-2025

```
import pandas as pd

from sklearn.datasets import load_diabetes

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler

from sklearn.ensemble import BaggingClassifier

from sklearn.tree import DecisionTreeClassifier

from sklearn.metrics import accuracy_score, confusion_matrix, mean_squared_error,
r2_score

import seaborn as sns

import matplotlib.pyplot as plt


# Load dataset

data = load_diabetes()

X = pd.DataFrame(data.data, columns=data.feature_names)


# Convert target to binary classification (1 = diabetic, 0 = not)

y = (data.target > data.target.mean()).astype(int)


# Split dataset

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


# Scale features

scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train)

X_test_scaled = scaler.transform(X_test)
```

```
# Train Bagging Classifier

model = BaggingClassifier(estimator=DecisionTreeClassifier(), n_estimators=50,
random_state=42)

model.fit(X_train_scaled, y_train)

# Predict

y_pred = model.predict(X_test_scaled)

# --- Metrics ---

accuracy = accuracy_score(y_test, y_pred)

mse = mean_squared_error(y_test, y_pred)

r2 = r2_score(y_test, y_pred)

cm = confusion_matrix(y_test, y_pred)

# --- Print Metrics ---

print("Accuracy:", round(accuracy, 4))

print("Mean Squared Error:", round(mse, 4))

print("R-squared Score:", round(r2, 4))

# --- Plot Confusion Matrix ---

sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Not Diabetic', 'Diabetic'],
            yticklabels=['Not Diabetic', 'Diabetic'])

plt.title("Confusion Matrix")


plt.xlabel("Predicted")

plt.ylabel("Actual")

plt.show()
```

OUTPUT:

model



accuracy


0.7303370786516854

mse

0.2696629213483146

r2

-0.0953846153846154



PROGRAM 17: Program to demonstrate Principal Component Analysis(PCA)

REGNO:24251115

DATE: 23-05-2025

```
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA

# Step 1: Load the Iris dataset
iris = load_iris()
X = iris.data          # Features (sepal & petal measurements)
y = iris.target        # Target (species)
target_names = iris.target_names

# Step 2: Standardize the data
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Step 3: Apply PCA to reduce dimensions (from 4 to 2)
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)

# Step 4: Create a DataFrame for PCA results
pca_df = pd.DataFrame(X_pca, columns=['PC1', 'PC2'])
pca_df['Species'] = y

# Step 5: Plot the 2 principal components
plt.figure(figsize=(7, 5))
for i, label in enumerate(target_names):
```

```
plt.scatter(pca_df[pca_df['Species'] == i]['PC1'],  
            pca_df[pca_df['Species'] == i]['PC2'],  
            label=label)
```

```
plt.title('PCA of Iris Dataset')
```

```
plt.xlabel('Principal Component 1')
```

```
plt.ylabel('Principal Component 2')
```

```
plt.legend()
```

```
plt.grid(True)
```

```
plt.show()
```


Program to demonstrate Principal Component Analysis(PCA)

OUTPUT:

```
scaler
```

```
StandardScaler
```

```
StandardScaler()
```



PROGRAM 18: .Program to demonstrate K Nearest Neighbours Classification(KNN).

REGNO:24251102

DATE: 23-05-2025

```
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, mean_squared_error
import seaborn as sns

# Load Iris dataset
iris = load_iris()
X = iris.data
y = iris.target
target_names = iris.target_names

# Split dataset
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Scale features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Train KNN model
knn = KNeighborsClassifier(n_neighbors=3)
```

```
knn.fit(X_train_scaled, y_train)

# Predict
y_pred = knn.predict(X_test_scaled)

# Evaluate
accuracy = accuracy_score(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)

print(f'Accuracy: {accuracy:.4f}')
print(f'Mean Squared Error: {mse:.4f}')

# Confusion matrix plot
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=target_names, yticklabels=target_names)
plt.title("Confusion Matrix - KNN")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```

OUTPUT:

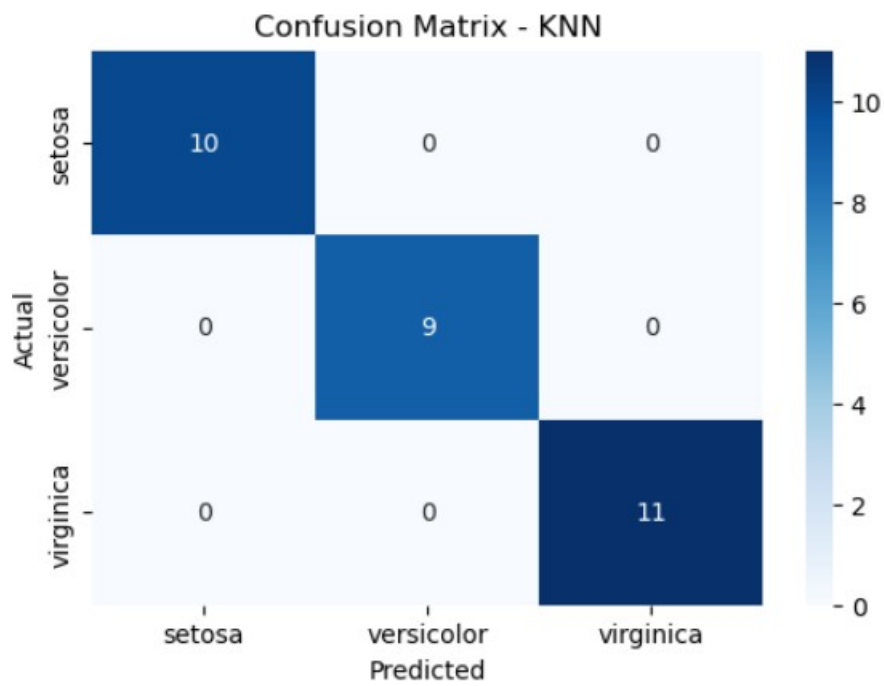
knn

KNeighborsClassifier

KNeighborsClassifier(n_neighbors=3)

Accuracy: 1.0000

Mean Squared Error: 0.0000



PROGRAM 19: Program to demonstrate Neural Networks and deep Learning

REGNO: 24251115

DATE: 23-05-2025

```
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import accuracy_score, confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

# 1. Load data
iris = load_iris()
X = iris.data
y = iris.target

# 2. Scale features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# 3. Split data
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2,
random_state=42)

# 4. Build and train the MLP neural network
mlp = MLPClassifier(hidden_layer_sizes=(10,), max_iter=500, random_state=42)
mlp.fit(X_train, y_train)
```

5. Predict

```
y_pred = mlp.predict(X_test)
```

6. Accuracy

```
acc = accuracy_score(y_test, y_pred)
```

```
print(f"Test accuracy: {acc:.4f}")
```

7. Confusion matrix

```
cm = confusion_matrix(y_test, y_pred)
```

```
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',  
             xticklabels=iris.target_names, yticklabels=iris.target_names)
```

```
plt.xlabel('Predicted')
```

```
plt.ylabel('Actual')
```

```
plt.title('Confusion Matrix')
```

```
plt.show()
```

OUTPUT:

scaler

StandardScaler 1 2
StandardScaler()


accuracy

1.0

y_pred

```
array([1, 0, 2, 1, 1, 0, 1, 2, 1, 1, 2, 0, 0, 0, 0, 1, 2, 1, 1, 2, 0, 2  
       0, 2, 2, 2, 2, 2, 0, 0])
```

	actual	predicted
0	1	1
1	0	0
2	2	2
3	1	1
4	1	1
5	0	0
6	1	1
7	2	2
8	1	1
9	1	1
10	2	2
11	0	0
12	0	0
13	0	0
14	0	0
15	1	1
16	2	2
17	1	1
18	1	1
19	2	2
20	0	0
21	2	2
22	0	0
23	2	2
24	2	2
25	2	2
26	2	2
27	2	2
28	0	0
29	0	0



PROGRAM 20: . Visualization using matplotlib for iris dataset

- **Scatter plot**
- **Box plot**
- **Line chart**

REGNO: 24251115

DATE: 23-05-2025

```
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
import pandas as pd

# Load Iris dataset
iris = load_iris()
df = pd.DataFrame(data=iris.data, columns=iris.feature_names)
df['species'] = pd.Categorical.from_codes(iris.target, iris.target_names)

# 1. Scatter plot: Sepal Length vs Sepal Width colored by species
plt.figure(figsize=(7,5))
for species in iris.target_names:
    subset = df[df['species'] == species]
    plt.scatter(subset['sepal length (cm)'], subset['sepal width (cm)'], label=species)
plt.xlabel('Sepal Length (cm)')
plt.ylabel('Sepal Width (cm)')
plt.title('Scatter Plot of Sepal Length vs Sepal Width')
plt.legend()
plt.show()

# 2. Box plot: Distribution of Petal Length by Species
plt.figure(figsize=(7,5))
```

```
df.boxplot(column='petal length (cm)', by='species')
```

```
plt.title('Box Plot of Petal Length by Species')
```

```
plt.suptitle("") # Removes default subtitle
```

```
plt.xlabel('Species')
```

```
plt.ylabel('Petal Length (cm)')
```

```
plt.show()
```

```
# 3. Line chart: Mean sepal length for each species
```

```
mean_sepal_length = df.groupby('species')['sepal length (cm)'].mean()
```

```
plt.figure(figsize=(7,5))
```

```
plt.plot(mean_sepal_length.index, mean_sepal_length.values, marker='o')
```

```
plt.title('Mean Sepal Length by Species')
```

```
plt.xlabel('Species')
```

```
plt.ylabel('Mean Sepal Length (cm)')
```

```
plt.grid(True)
```

```
plt.show()
```

OUTPUT:

